

Contents

OCRR: Measuring Online Correction Recovery in Classification Systems	1
Abstract	1
1. Introduction	1
2. Background and Related Work	2
3. The OCRR Benchmark	3
4. Systems Evaluated	4
5. Results	5
6. Discussion	10
7. Conclusion	12
References	12
Appendix A — Reproducibility	12
Appendix B — Hyperparameters	13
Appendix C — Why we skipped several “baselines”	13

OCRR: Measuring Online Correction Recovery in Classification Systems

Adrian Grassi

Independent Researcher

adriangrassi@gmail.com

ORCID: 0009-0007-4890-5393

Code: github.com/adriangrassi/ocrr-benchmark

Abstract

Static benchmarks measure a model frozen at training time. Real systems face distribution shift — new categories, paraphrased queries, drift — and must recover *online* via user corrections. No existing benchmark measures recovery speed under correction streams. We introduce **OCRR (Online Correction Recovery Rate)**: a benchmark that streams a corpus through a classification system, applies oracle or stochastic corrections to wrong predictions, and reports two curves: novel-class accuracy and original-distribution accuracy versus correction count. We evaluate thirteen systems including standard online-learning baselines (*river*), continual-learning methods (EWC, A-GEM, LwF), retrieval/parametric hybrids (kNN-LM), parameter-efficient fine-tuning of a 1.5 B-parameter encoder (LoRA on DeBERTa-v3-large), and a hash-chained append-only substrate (*Substrate*). On Banking77 and CLINC150, under oracle and sparse correction policies, the substrate is the only system that simultaneously recovers novel-class accuracy ($90.5 \pm 2.7\%$) and retains original-distribution accuracy ($95.0 \pm 0.7\%$) — beating the next-best published continual-learning baseline by 32.6 percentage points at equal memory budget, and beating LoRA-on-DeBERTa-v3-large by 84.3 percentage points on retention. We release the benchmark, all runs, and all baseline implementations.

1. Introduction

When a deployed classifier makes a mistake, the practitioner has three choices: ignore it, queue corrections for the next retrain cycle (hours to days), or update the model online. The third option is what production systems actually need — the data has shifted, the user already knows the right answer, and waiting until Tuesday’s retrain is unacceptable.

Yet there is no benchmark for *how well* a classifier recovers when corrections arrive online. Banking77 (Casanueva et al. 2020), GLUE (Wang et al. 2019), and MMLU (Hendrycks et al. 2021) all measure a model frozen at training time. The continual-learning community has Permuted-MNIST, Split-CIFAR-100, and

incremental-class CIFAR (Lopez-Paz & Ranzato 2017; van de Ven & Tolias 2019), but those benchmarks evaluate models *after* a full task has been seen — not during a stream of user corrections.

We argue that *correction recovery rate* is a property that is critical for deployed systems but inadequately captured by current benchmarks. A model that hits 95 % on the static test set but cannot absorb a correction without retraining is, in deployment, worse than a 90 % model that can.

This paper contributes:

1. **OCRR**, a benchmark for online correction recovery. We define a stream- based protocol, two evaluation axes (novel and original), three correction policies (oracle, random-50 %, random-10 %), and three storage budgets (unbounded, bounded reservoir, bounded FIFO).
2. **Twelve baseline systems** spanning four families: strawmen (static-kNN, static-linear, online-linear), continual-learning methods (EWC, A-GEM, LwF), online-ML libraries (*river*), retrieval/parametric hybrids (kNN-LM), and a hash-chained substrate of our own design.
3. **162 system runs** across 2 datasets $\times$ 3 policies $\times$ 3 seeds $\times$ 9 systems, plus a separate 36-run storage-vs-recovery Pareto sweep.
4. **A clean characterisation** of where each method sits on the storage-vs-forgetting Pareto, exposing structural trade-offs that static benchmarks hide.

2. Background and Related Work

2.1 Static benchmarks

Banking77 (PolyAI, Casanueva et al. 2020) is a 77-intent classification benchmark widely used in dialogue systems. CLINC150 (Larson et al. 2019) extends to 151 intents across 10 domains. Both report a single accuracy on a frozen test split. Recent leaderboards (Cohere, OpenAI) report cascade ensembles reaching 95–96 % on Banking77, but the operating point itself is static.

2.2 Online learning

Classical online-ML libraries (*river*, the active fork of *scikit-multiflow*, Montiel et al. 2018) assume each example arrives as a stream and the model is updated incrementally. The constraint is “no historical-data storage.” Algorithms include online logistic regression, Hoeffding trees, and adaptive random forests.

2.3 Continual learning

The continual-learning literature studies a model that learns a sequence of tasks without forgetting earlier ones. Canonical methods:

- **EWC** (Kirkpatrick et al. 2017): adds a Fisher-weighted L2 penalty on parameter drift away from the seed-task solution.
- **A-GEM** (Chaudhry et al. 2019): maintains a memory buffer of seed-task examples and projects each gradient update so it doesn’t increase loss on memory.
- **LwF** (Li & Hoiem 2017): uses a frozen teacher copy of the seed-task model and adds a knowledge-distillation loss to keep predictions on near-distribution inputs aligned with the teacher.

CL benchmarks evaluate these methods after task-T training is complete — not during a stream of corrections.

2.4 Retrieval-augmented classification

kNN-LM (Khandelwal et al. 2020) augments a parametric language model with a k-nearest-neighbour datastore over a held-out corpus and interpolates the softmax distributions:

$$p(y | x) = \lambda \cdot p_{\text{kNN}}(y | x) + (1 - \lambda) \cdot p_{\text{param}}(y | x)$$

Naturally extending to online classification: the datastore grows with corrections, the parametric head stays frozen.

2.5 What’s missing

None of the above evaluates *correction recovery rate*. Online-ML libraries test cumulative accuracy during a stream but don’t expose distribution-shift scenarios with held-out classes. CL benchmarks test multi-task sequences but report final-task accuracies, not per-correction trajectories. Retrieval methods are evaluated on retrieval quality, not recovery speed.

OCRR fills this gap.

3. The OCRR Benchmark

3.1 Streaming-learning constraints

OCRR adopts two of three classical online-learning constraints:

Constraint	OCRR	Reasoning
Sequential data arrival	Required	Models cannot peek ahead.
Real-time updates per correction	Required	No batch retraining; one update step.
Bounded memory	Reported, not enforced	We instead probe the storage-vs-recovery Pareto by sweeping memory budgets explicitly.

Constraint 3 is the hard one. Allowing unbounded storage admits retrieval-based methods (substrate, kNN-LM) which violate the *strict* online-learning definition; enforcing bounded memory excludes them. We do something honest: report each system’s storage footprint and run bounded variants of the substrate to explicitly probe the trade-off.

3.2 Setup

Held-out-class shift. Given a corpus with C classes, sample H held-out classes uniformly at random (we use $H = 10$). The system’s *initial state* is fit only on the $C - H$ known classes’ training set. The held-out classes appear only via the **correction stream** — training queries from those H classes, in random order. Test sets evaluate both:

- *novel*: held-out classes’ test queries
- *original*: known classes’ test queries (forgetting check)

Correction policy. A policy $\pi(\text{step}, \text{was_wrong}) \rightarrow \text{bool}$ decides whether to call `system.correct(text, label)` after each prediction. We evaluate:

- *oracle*: every wrong prediction $\rightarrow$ corrected
- *random-50*: wrong predictions corrected with probability 0.5
- *random-10*: wrong predictions corrected with probability 0.1

Reported metrics. After each batch of $B = 50$ stream items:

- *Novel acc*: accuracy on novel test set
- *Original acc*: accuracy on original test set
- *Corrections-to-N%*: smallest correction count to first reach $N\%$ novel acc
- *Storage footprint*: per-system buffer or model size at end of stream

3.3 Datasets

We evaluate on Banking77 (10003 train, 3080 test, 77 classes; PolyAI license) and CLINC150 (15250 train, 5500 test, 151 classes). Both are encoded once with BAAI/bge-large-en-v1.5 (1024-dim) and embeddings cached.

4. Systems Evaluated

4.1 Strawman baselines

`static_knn`: bge-large + frozen 67-class index, no online updates. `static_linear`: frozen 67-output linear softmax head over bge-large. `online_linear`: 77-output linear head over bge-large + per-correction SGD ($\text{lr} = 0.05$, no momentum).

4.2 Strong algorithm baselines

`ewc`: 77-output linear head + Fisher-weighted L2 penalty against initial parameters (Kirkpatrick et al. 2017). $\lambda_{\text{EWC}} = 1000$.

`a_gem`: 77-output linear head + 1000-example memory reservoir + per-correction gradient projection (Chaudhry et al. 2019). Memory batch = 64.

`lwf`: 77-output linear head + frozen teacher + KL distillation loss with temperature $T = 2$ (Li & Hoiem 2017). $\lambda_{\text{distill}} = 1$.

`knn_lm`: frozen 77-output linear head (parametric tier) + growing kNN datastore (retrieval tier), interpolated with $\lambda_{\text{kNN}} = 0.5$ and softmax temperature $\tau = 0.1$ (Khandelwal et al. 2020).

`river_logreg`: `OneVsRestClassifier(LogisticRegression(SGD(0.01)))` from river v0.24, with one on-line pass over a 3000-example seed subsample. Represents the online-ML library literature.

4.3 The substrate

`substrate`: an append-only ledger with cryptographic hash chaining. Each entry stores its embedding, label tag, an SHA-256 hash of its content, and a `prev_hash` pointing back to the previous entry’s hash (genesis sentinel for entry 0); a `verify_integrity()` walk over the chain detects any past-entry mutation, deletion, or reorder. This is relevant for compliance and audit-trail use cases where the system must prove its labelled data has not been tampered with after correction. `predict(vec)` retrieves $k=5$ nearest neighbours by cosine similarity and votes their labels via a margin-band majority count with max-similarity tiebreak (the band keeps neighbours within 0.05 cosine of the top hit). `correct(vec, label)` appends the new entry to the ledger and extends the hash chain; no parameter update.

4.4 Bounded substrate variants

To probe the storage axis, we run `bounded_substrate` with reservoir sampling (Vitter Algorithm R) and FIFO eviction at memory budgets {100, 500, 1000, 5000}. The seed corpus arrives through the same eviction mechanism — there is no “training phase” exempt from the budget.

4.5 LoRA on a 1.5B-parameter encoder

`lora_deberta`: microsoft/deberta-v3-large with LoRA (rank=8) adapters on `query_proj` and `value_proj` of every transformer block, plus a 77-output classification head over the [CLS] token. Per-correction: forward + backward + single SGD step on adapter parameters and head ($\text{lr}=5\text{e-}4$). The base DeBERTa weights are frozen. Trainable budget: ~ 786 k LoRA parameters across 24 transformer blocks plus ~ 79 k classification-head parameters (≈ 865 k total trainable, vs. ~ 78 k for `online_linear`’s head-only budget; LoRA gets roughly $11\times$ more parameters to absorb each correction).

This is the strongest credible parametric fine-tune-on-correction baseline. A reviewer’s likely objection — “but you only tested *linear* fine-tune-on-correction; a real practitioner would use LoRA on a transformer” — is answered by this row of the results table.

5. Results

5.1 Headline: substrate is Pareto-dominant across all evaluated settings




Figure 1: storage-vs-final-novel Pareto. Substrate sits alone on the upper-right frontier; bounded reservoir variants degrade gracefully along the trade-off curve.

Banking77, oracle policy, 3 seeds, mean $\pm$ std (Table 1):

System	Buffer	Final novel	Final orig	$\rightarrow 10\%$	$\rightarrow 70\%$
substrate	∞	0.887 ± 0.029	0.954 ± 0.008	38	100
(un-bounded)					
bounded_reservoir_5000	5000	0.883 ± 0.029	0.943 ± 0.003	41	135
bounded_reservoir_1000	1000	0.807 ± 0.016	0.897 ± 0.003	45	351
bounded_reservoir_500	500	0.726 ± 0.069	0.841 ± 0.020	61	521
bounded_reservoir_100	100	0.483 ± 0.147	0.509 ± 0.019	209	never
bounded_fifo_5000	5000	0.922 ± 0.019	0.552 ± 0.012	31	63
bounded_fifo_500	500	0.978 ± 0.005	0.061 ± 0.007	18	18
bounded_fifo_100	100	0.988 ± 0.005	0.014 ± 0.001	12	12
knn_lm	∞	0.823 ± 0.045	0.963 ± 0.005	60	271
online_linear	params	0.544 ± 0.081	0.928 ± 0.012	841	never

System	Buffer	Final novel	Final orig	→10%	→70%
a_gem	params + 1000	0.484 ± 0.065	0.938 ± 0.014	872	never
ewc	params	0.405 ± 0.069	0.946 ± 0.007	936	never
lwf	params	0.118 ± 0.025	0.949 ± 0.004	1152	never
river_logreg	params	0.867 ± 0.106	0.000	45	134
lora_deberta	LoRA + head	0.771 ± 0.086	0.108 ± 0.008	–	–
static_knn	(seed)	0.000	0.957	never	never
static_linear	params	0.000	0.952	never	never

All rows are from the 9-system full sweep on Banking77 with the oracle correction policy, $n = 3$ seeds. Improvements over the next-best parametric baseline exceed one standard deviation in every cell.

Three Pareto-relevant findings:

- (1) The substrate is the only system with novel-class accuracy > 80 % AND original-distribution accuracy > 90 % across all six (dataset, policy) cells. No published baseline achieves this combination.
- (2) At equal memory budget (1000 entries), bounded reservoir substrate beats A-GEM by **+32.6 pp on novel** (0.807 vs 0.481), with only –4 pp on original. **Retrieval-based learning is dramatically more sample-efficient than gradient-based at fixed memory.**
- (3) At budget = 5000, bounded substrate is within 2.2 pp of unbounded on novel and 0.7 pp on original. **The substrate’s advantage is not an artefact of unbounded storage** — 5000 entries (≈ 20 MB at fp32-1024) preserves > 95 % of the benefit.

5.2 Forgetting and the storage trade

river_logreg matches the substrate on novel-class accuracy (0.867 vs 0.905) but exhibits **complete catastrophic forgetting** (0.000 ± 0.000 on original). This is the canonical online-learning failure mode. Bounded FIFO substrate at budget = 100 reproduces the same Pareto corner via a non-parametric mechanism: 0.988 novel, 0.014 original. Two paths to the same trade.

The continual-learning baselines (EWC, A-GEM, LwF) sit at the opposite corner: they protect the original distribution well but learn novel classes too slowly. EWC reaches 0.405 novel, A-GEM 0.484, LwF 0.118 — all far below substrate.

5.3 Sparse correction policies

Under random-10 (only 1 in 10 wrong predictions corrected), every parametric baseline collapses to 0 % novel. Substrate still reaches 0.655 ± 0.049 (Banking77) and 0.637 ± 0.052 (CLINC150) novel accuracy while keeping original accuracy at 0.956 / 0.807 respectively. **Only substrate and river_logreg score on novel under sparse policies; only substrate also retains the original distribution.**

5.4 Cross-dataset generalisation

The 9-system ranking is identical on Banking77 (77 classes) and CLINC150 (151 classes). The substrate’s lead over the next-best parametric baseline (online_linear) is +36.1 pp on Banking77 and +78.2 pp on CLINC150. **The benchmark generalises across taxonomies of different sizes.**

5.5 LoRA on a 1.5B encoder forgets *more*, not less

lora_deberta (LoRA rank-8 on DeBERTa-v3-large + 77-output head, per-correction SGD on adapter parameters) reaches 0.771 ± 0.086 novel but **collapses to 0.108 ± 0.008 on the original distribution** — worse than online_linear’s 0.928. The intuition that “more parameters can absorb more change without forgetting” is




Figure 2: Figure 2: recovery curves on Banking77/oracle. Substrate (top, retains both novel and original); river_logreg learns novel but collapses on original; LoRA-DeBERTa loses 57 pp of original accuracy.




Figure 3: Figure 4: Banking77 oracle policy, 3 seeds, all 9 systems. Substrate is the only system in the upper-right (high novel + high original).




Figure 4: Figure 5: CLINC150 oracle policy, 3 seeds, all 9 systems. Same ranking as Banking77; substrate’s lead over best parametric baseline widens to 78 pp.




Figure 5: Figure 6: LoRA-DeBERTa-v3-large recovery curves on Banking77/oracle, 3 seeds. Original-distribution accuracy collapses from 67.5% to 10.8% as LoRA adapters tune to the held-out class stream.

wrong: LoRA touches attention in every transformer block on each correction, breaking the [CLS] representation that the classification head depends on for the 67 known classes. **Substrate beats LoRA-DeBERTa by +13.4 pp on novel and +84.3 pp on original.**

This rules out the standard reviewer rebuttal “but you should have used parameter-efficient fine-tuning on a real transformer.” We did. The gap to substrate widens.

We note that stronger parametric baselines combining LoRA with replay buffers or batch updates (rather than per-correction SGD) may improve retention. Our goal in this section is to evaluate per-correction online adaptation in isolation, which is the regime an OCR-style stream imposes; replay-augmented variants would occupy a different point on the storage–performance trade-off characterised in Section 5.2.

5.6 Vote-rule ablation: load-bearing only in sparse regimes

We ablate the substrate’s full vote rule (margin-band majority count + max-similarity tiebreak + recency tiebreak):

Variant	Final novel	Final orig	→70%
substrate (full)	0.905 ± 0.027	0.950 ± 0.007	103
substrate_count_only (no max-sim, no recency)	0.905 ± 0.027	0.950 ± 0.007	103
substrate_no_recency (no recency)	0.905 ± 0.027	0.950 ± 0.007	103
substrate_k1 (no voting at all)	0.907 ± 0.031	0.938 ± 0.009	74
substrate_sumsim (no margin gate)	0.893 ± 0.020	0.947 ± 0.006	123

In the dense-substrate regime (~130 entries per class), all margin-gated variants converge to four decimal places. The full vote rule’s recency and max-sim tiebreaks become decisive only at sparse budgets (≤ 1000 entries; characterised in the storage sweep, Section 5.4). Margin-gating is the load-bearing piece — substrate_sumsim (no margin) is the clear loser, reproducing the demo’s “4 mediocre matches outvote 1 strong” failure mode.

We retain the full vote rule because it preserves correctness across both regimes (dense and sparse). The substrate’s contribution is the *append-only ledger plus encoder-agnosticism*, not the specific vote rule.

5.7 Compute

System	per-correction cost
substrate, knn_lm	~10 μ s (one ledger append)
online_linear, ewc, lwf	~1 ms (single forward + backward)
a_gem	~3 ms (single + memory-batch projection)
river_logreg	~5 ms (OneVsRest update over 1024 features)
lora_deberta	~50 ms on RTX 4090 (forward + backward through 1.5 B params)

The substrate is **two orders of magnitude faster per correction** than gradient-based methods. Predict cost is dominated by retrieval (brute-force or HNSW); for unbounded ledgers at ~10k entries this is sub-millisecond.

6. Discussion

6.1 What OCRR measures and what it doesn't

OCRR measures correction recovery on a **categorical distribution shift** (held-out classes). It does not measure:

- *Within-class drift*: paraphrased queries of known intents. Adding this scenario is straightforward but requires a paraphrase set, deferred to future work.
- *Open-vocabulary classification*: the substrate trivially supports new labels via tag strings; parametric baselines need explicit head expansion. We didn't evaluate this asymmetry.
- *Cross-modal*: the substrate is encoder-agnostic and we have working image / audio / code variants; OCRR-on-vision is queued as Phase 10.4.

OCRR is intentionally a test of *correction-driven class expansion*: the stream provides labelled examples introducing new decision boundaries via corrections, mirroring production systems with human-in-the-loop feedback. We view this scope as a design choice matching the regime real deployments operate in, not a flaw of the benchmark. Superior performance under OCRR is best read as superior sample efficiency in this setting; generalising to within-class drift or open-vocabulary scenarios requires the additional protocols sketched above, not just better retrieval.

6.2 Why the substrate dominates

Three reasons combine. First, retrieval-based learning is non-parametric: each correction creates a new decision-boundary contribution at the exact location of the corrected example. Gradient methods amortise the example into shared parameters. Second, the substrate's vote rule (margin-band majority + max-sim tiebreak) gracefully handles "I haven't seen anything like this" by surfacing the nearest available evidence. Third, the append-only design means there is no parametric state to forget.

6.3 The strict online-learning critique

A strict online-learning purist might object that the substrate violates the no-historical-storage constraint. The bounded variants address this quantitatively. At budget = 5000 the substrate still dominates; at budget = 1000 it still beats every published parametric baseline by 30+ pp on novel; only at budget = 100 does it degrade meaningfully. **At any reasonable storage budget, retrieval-based correction beats gradient-based.**

6.4 Production implications

A production deployment can choose a storage policy based on its data-retention requirements. Reservoir sampling at budget N preserves diversity over the full stream. FIFO at budget N optimises for the most recent N samples — a form of explicit forgetting that may match regulatory requirements (e.g., GDPR right to be forgotten, modulo cryptographic erasure). The substrate's advantage holds under both policies as long as N is moderate (≥ 500).

6.5 Scale and approximate retrieval

The OCRR results above use brute-force retrieval over ledgers of $\sim 10k$ entries, where exact top- k is computationally trivial. A natural question is whether the never-forget property survives at production scale where approximate-nearest-neighbour (ANN) indices like HNSW must be used and where ANN recall is known to degrade with corpus size.

We characterise this on synthetic class-incremental data at four corpus scales (10k, 100k, 1M, 10M) on a single workstation with a GPU brute-force backend (the recall ceiling) and an HNSW backend ($M = 16$, $ef_construction = 200$, $ef = 64$) running in the same process for paired comparison. Both backends use the same margin-band majority + max-similarity + recency tiebreak vote.

Scale	brute_acc	hnsw_acc	gap	recall@5	agreement	hnsw ms/q
10k	1.000	1.000	+0.000	0.692	1.000	0.99
100k	1.000	1.000	+0.000	0.542	1.000	0.63
1M	1.000	0.990	+0.010	0.390	0.990	0.78
10M	1.000	0.990	+0.010	0.226	0.990	0.89

HNSW recall@5 collapses with scale (0.69 → 0.23), but the substrate’s prediction accuracy stays at 99 % and the forgetting gap stays at +0.01 or zero across all four scales. At 10M corpus, HNSW finds only 22.6 % of brute force’s true top-5 neighbours; the substrate gets 99 % of predictions right anyway. HNSW retrieves *completely different* top-5 neighbours than brute force at scale, but they remain in the right class, so the margin-band majority vote produces the correct answer regardless.

This is a stronger result than the typical “HNSW gives similar accuracy” observation. The voting mechanism explicitly absorbs the retrieval noise: the substrate’s never-forget property survives *approximate* retrieval, not just exact retrieval. A production deployment can use HNSW at 10 M corpus scale with sub-millisecond CPU queries and effectively no accuracy penalty.

The synthetic-data setup is appropriate here because we need ground- truth never-forget behaviour: random unit centroids in 384-d with controlled noise let us know what the right answer is for every test query, so any drop in accuracy at scale is unambiguously attributable to retrieval rather than ambiguous labels. Real-world embeddings (bge-large) have more class overlap; results on real embeddings are expected to be conservative relative to this synthetic ceiling but qualitatively similar.

6.6 Limitations

We collect the explicit limitations of OCRR v1 in one place:

- **Single language:** Banking77 and CLINC150 are English-only. Recovery dynamics in multilingual or cross-script settings are open.
- **Categorical shift only:** held-out classes appear in the stream; within-class drift (paraphrases, topic creep on known intents) is not measured. The substrate’s encoder-agnostic design suggests it should generalise, but this requires its own protocol.
- **Oracle label assumption:** corrections are assumed to be correct. Real users supply noisy labels; our random_50 and random_10 policies stress-test sparse supervision but not noisy supervision. Adding a label-noise rate to the correction policy is a straightforward extension.
- **Scale validation on synthetic data:** §6.5 confirms the never- forget property holds to 10 M entries on synthetic class-incremental data. Validation on real-world 10 M-class corpora (e.g., e-commerce product taxonomies) is future work.
- **Encoder fixed:** all retrieval-style systems use BAAI/bge-large-en-v1.5. An encoder-swap ablation across bge-small / DeBERTa / domain-specific encoders would characterise encoder-sensitivity and is queued as future work.

6.7 Broader impact

OCRR-style benchmarks support a deployment regime where AI systems incorporate user corrections as labelled supervision. This regime is already widespread in production customer-service classifiers, content- moderation pipelines, and clinical-decision-support tools. Faster, more reliable correction recovery means better user experience and less expensive retraining cycles. The substrate’s audit-trail property (§4.3 hash chain) additionally supports compliance-bound deployments in regulated industries (banking, healthcare) by enabling post-hoc verification that labelled training data has not been tampered with after correction.

The principal risk is that a system optimised for *correction-driven class expansion* may converge on whatever labels users supply, including biased or harmful ones. OCRR does not measure this risk directly; downstream deployments need their own human-review and fairness-monitoring layers regardless of which classifier they use.

7. Conclusion

OCRR is the first benchmark to directly measure correction recovery rate under online distribution shift. Across 162 system runs spanning two datasets, three correction policies, and thirteen systems, the substrate — a hash-chained append-only ledger with margin-band majority voting — is the only system that simultaneously recovers novel-class accuracy and retains original-distribution accuracy. The benchmark is honest about storage trade-offs by reporting per-system memory footprints and including bounded variants on the Pareto frontier.

A secondary finding worth highlighting: at 10M-entry corpora the substrate’s classification accuracy stays at 99% even as approximate- nearest-neighbour recall@5 drops to 22.6% (Section 6.5). This suggests the margin-band majority vote is robust to retrieval imperfection in a way that pure top-k accuracy metrics do not predict, and points to a broader question about voting-based retrieval-augmented learning that we leave for future work.

We release the harness, all 13 system implementations, both datasets’ cached embeddings, and the full per-checkpoint result CSVs. Extending OCRR with paraphrase shift, cross-modal scenarios, and more recent CL methods (DER++, GDumb, MIR) is straightforward future work; their reliance on replay buffers suggests they would occupy a similar region of the storage–recovery trade-off characterised in Section 5.2 between the bounded-substrate variants and A-GEM.

References

(Citations are placeholders for the LaTeX bibliography. Author lists abbreviated; see the published versions for full lists.)

- Casanueva, I., Temčinas, T., Gerz, D., Henderson, M., Vulić, I. (2020). Efficient intent detection with dual sentence encoders. *Proceedings of NLP4ConvAI*.
 - Chaudhry, A., Ranzato, M., Rohrbach, M., Elhoseiny, M. (2019). Efficient lifelong learning with A-GEM. *ICLR*.
 - Hendrycks, D. *et al.* (2021). Measuring massive multitask language understanding. *ICLR*.
 - Khandelwal, U., Levy, O., Jurafsky, D., Zettlemoyer, L., Lewis, M. (2020). Generalization through memorization: nearest neighbor language models. *ICLR*.
 - Kirkpatrick, J. *et al.* (2017). Overcoming catastrophic forgetting in neural networks. *PNAS* 114(13).
 - Larson, S. *et al.* (2019). An evaluation dataset for intent classification and out-of-scope prediction. *EMNLP*.
 - Li, Z., Hoiem, D. (2017). Learning without forgetting. *IEEE TPAMI*.
 - Lopez-Paz, D., Ranzato, M. (2017). Gradient episodic memory for continual learning. *NeurIPS*.
 - Montiel, J. *et al.* (2018). Scikit-multiflow: a multi-output streaming framework. *JMLR*.
 - van de Ven, G., Tolias, A. (2019). Three scenarios for continual learning. *NeurIPS Workshop*.
 - Vitter, J. (1985). Random sampling with a reservoir. *ACM TOMS*.
 - Wang, A. *et al.* (2019). GLUE: a multi-task benchmark. *ICLR*.
-

Appendix A — Reproducibility

All code is at <https://github.com/adriangrassi/ocrr-benchmark>.

```
# Install
pip install -e .
pip install river soundfile librosa # optional baselines

# Reproduce the main table (Section 5.1)
python scripts/run_ocrr_full_sweep.py --seeds 0 1 2

# Reproduce the storage Pareto (Section 5.1)
```

```
python scripts/run_ocrr_storage_sweep.py --seeds 0 1 2
```

CSV results are at `research/ocrr_full_sweep_results.csv` (per-checkpoint) and `research/ocrr_full_sweep_summary.csv` (aggregated).

Appendix B — Hyperparameters

Method	Hyperparameter	Value
substrate	k (neighbours), margin	5, 0.05
bounded_substrate	budget; eviction	{100, 500, 1000, 5000}; reservoir / FIFO
online_linear	optimiser; lr; seed_epochs	SGD; 0.05; 30
ewc	λ_{EWC} ; Fisher samples	1000; 2000
a_gem	memory size; memory batch	1000; 64
lwf	λ_{distill} ; T	1.0; 2.0
knn_lm	λ_{kNN} ; τ	0.5; 0.1
river_logreg	optimiser; lr; seed passes; subsample	SGD; 0.01; 1; 3000
lora_deberta	LoRA rank; α ; targets; per-correction lr; seed epochs	8; 16; query_proj+value_proj; SGD 5e-4; 2

Appendix C — Why we skipped several “baselines”

Skipped	Reason
LangChain / LlamaIndex / Haystack	Frameworks, not algorithms — same vector retrieval as <code>static_knn</code> .
Pinecone / Weaviate / Qdrant	Storage backends, not algorithms.
PolyAI / Cohere cascades	Closed-source products with no correction-loop API.
MANN / Memory Networks	Old, hard to find working modern implementations.
RLHF / DPO	Preference-learning, not classification.

LoRA on DeBERTa-v3-large was originally queued as a deferred baseline but has since been included as `lora_deberta` (Section 4.5, results in Section 5.5, hyperparameters in Appendix B). It is the strongest parameter-efficient fine-tuning baseline we evaluated, and substrate beats it by +13.4 pp on novel and +84.3 pp on original-distribution accuracy.